

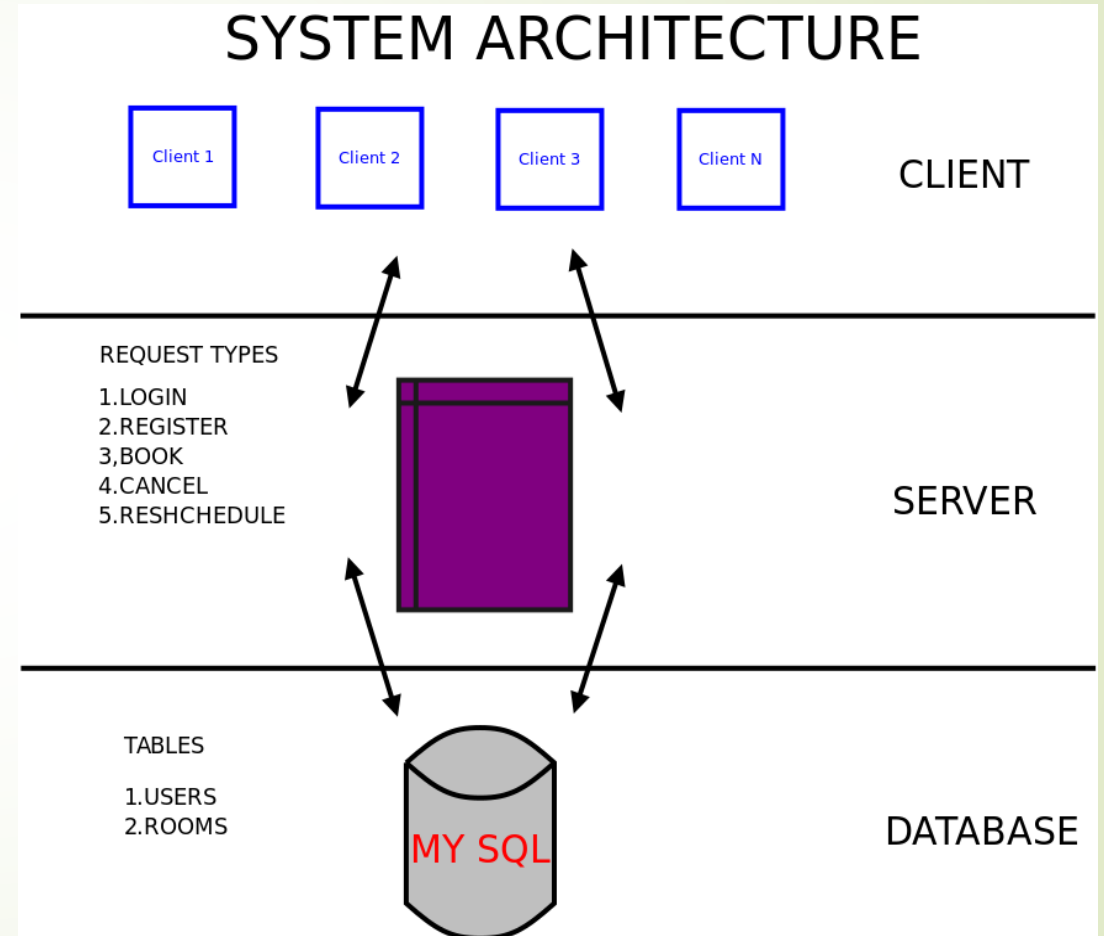


Room Booking System

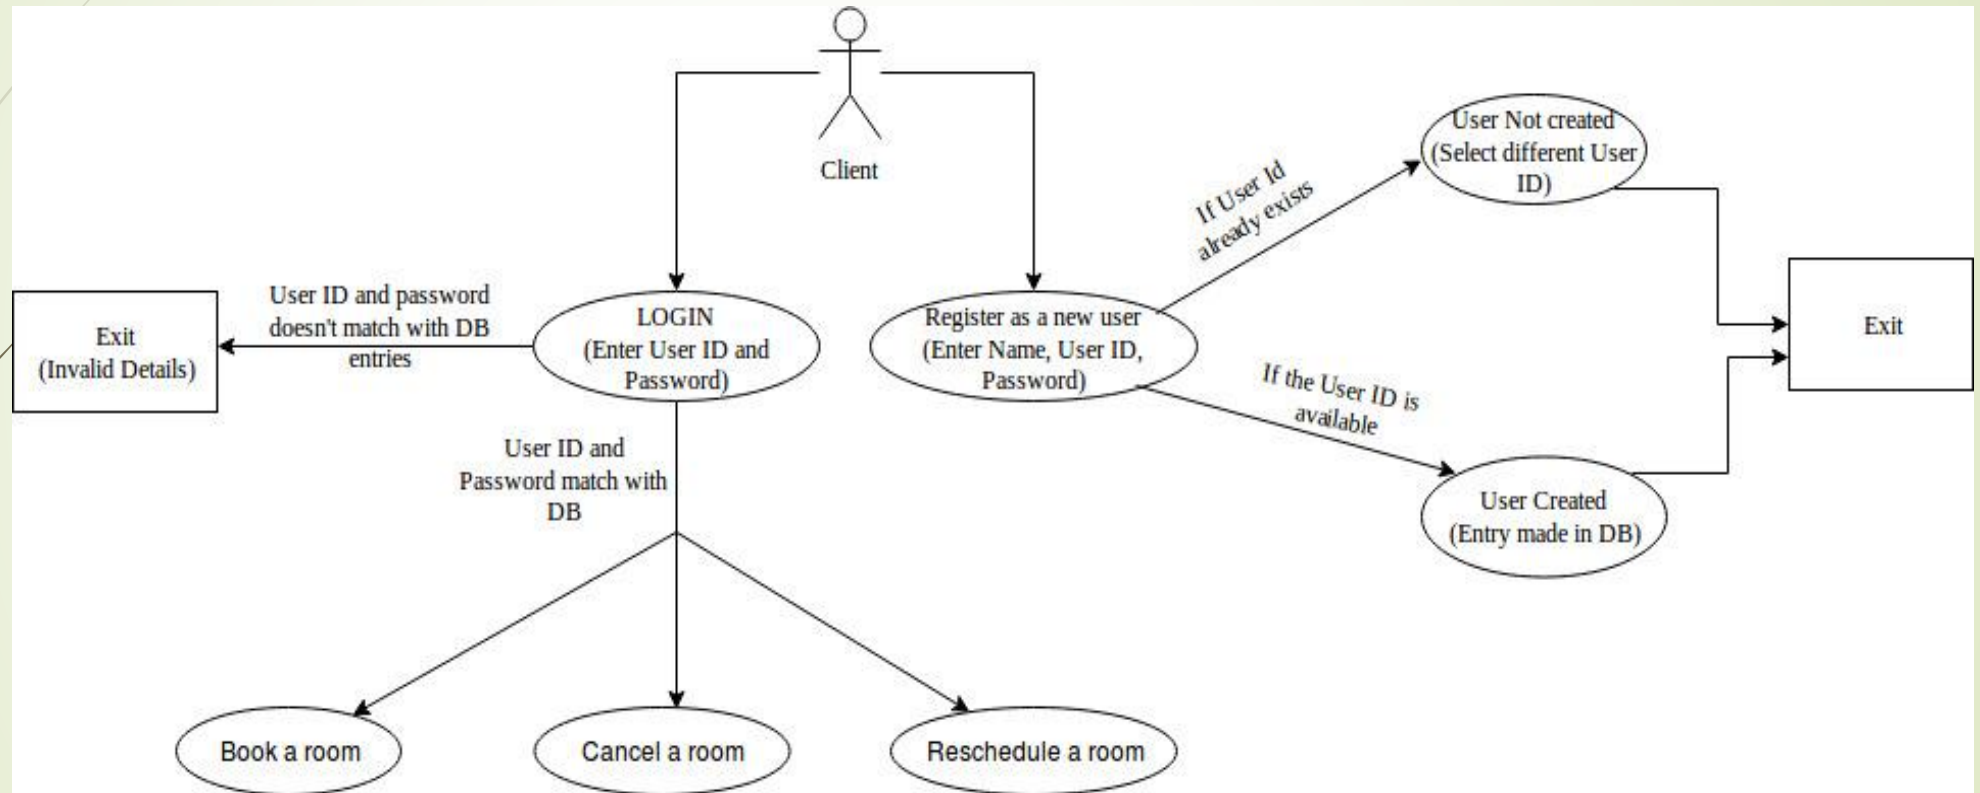
Bhagyesh Patil (16305R003), Ayush Goyal (16305R011)

System Architecture

- Multi-threaded.
- Front-end server:
 - Clients interact with it.
 - Processes their requests.
 - Interact with backend server.
- Back-end server:
 - Stores the information about users and rooms.



Flow Diagram





Phase 1: Constructing the system

- Database used: MySQL.
- Types of requests: Register, Login, Book, Cancel, Reschedule.
- Critical Sections: Booking, rescheduling.
- User authentication required to enter the system.
- Booking and rescheduling successful if room is available on required date.
- Date check.
- Unique username.
- Room number given at the time of booking is used for Canceling and Rescheduling.


Phase 2: Load Testing

- Closed Loop testing. No think time between requests.
- Increased the load on the server by increasing the number of concurrent users.
- The load generator can be used to perform load testing of:
 - Book requests only.
 - Mixture of Book and Cancel requests.
- Identified bottleneck resource using htop and system monitor.
- **Changed the configuration file of the MySQL.**
- **Reduced the unnecessary locking of some instructions in the critical section.**
- Bottleneck: CPU consumption by MySQL.
- The system hits saturation when the number of users are 300 (for both requests).
- The capacity of the system when it hits bottleneck is around 415 requests/sec (for both requests).


Obtained Graphs in Phase 2

- Book Request Only: (Run time of experiment = 4 minutes)

Average Throughput vs Load




Avg Response Time vs Load




- Mixture of Book (70%) and Cancel Requests(30%):

Average Throughput vs Load



Average Response Time vs Load





Phase 3: Optimizing System Performance



- First approach:
 - Moved the MySQL database to In-memory.
 - Increased the performance (less as compared to Redis).
- Second approach:
 - Used **Redis** (In-memory DB, stores in key-value pair).
 - Data read and written to RAM.
 - User requests are fulfilled faster as compared with MySQL, so we get high throughput and low response time.
 - Used redis hashes to store each tuple.
 - Bottleneck: CPU consumption by Front-end server.
 - The system hits saturation when the number of users are 25 (for both requests).
 - The capacity of the system is around 7700 req/sec (for book only) and 9450 req/sec (for mixture of book and cancel).

Obtained Graphs in Phase 3

- Book Request Only: (Run time of experiment = 4 minutes)



- Mixture of Book (70%) and Cancel Requests(30%):



Comparison between Phase 2 and 3

➤ **The throughput is increased drastically in Phase 3**

- As Redis handles requests directly from RAM which takes less time compared to insertion and removal from disk.
- RAM is fast for I/O and disks are comparatively slow.
- System hits saturation with less number of users.

➤ **The response time is decreased drastically in Phase 3**

- As the number of requests being fulfilled are increased by a large number as insertion is done in RAM which takes very less time.
- In case of MySQL response time was hundreds of milliseconds while in redis response time is hundreds of microseconds.

➤ **Understandings:**

- Try to use minimal locking as it decreases performance drastically.



Thank You.